

AI MAJOR AUGUST BATCH

Project Name:EDA ON SALARY DATASET

```
[2] import pandas as pd
df = pd.read_csv('https://github.com/ameermanna8824/DATASETS/raw/main/Salary_Data.csv')
df
```

	YearsExperience	Salary
0	1.1	39343.0
1	1.3	46205.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0
5	2.9	56642.0
6	3.0	60150.0
7	3.2	54445.0
8	3.2	64445.0
9	3.7	57189.0
10	3.9	63218.0
11	4.0	55794.0
12	4.0	56957.0
13	4.1	57081.0
14	4.5	61111.0

```
✓ [2] 15      4.9  67938.0
      16      5.1  66029.0
      17      5.3  83088.0
      18      5.9  81363.0
      19      6.0  93940.0
      20      6.8  91738.0
      21      7.1  98273.0
      22      7.9 101302.0
      23      8.2 113812.0
      24      8.7 109431.0
      25      9.0 105582.0
      26      9.5 116969.0
      27      9.6 112635.0
      28     10.3 122391.0
      29     10.5 121872.0
```

```
✓ [4] #years of experience and salary
      x = df.iloc[:,0:1].values
      y = df.iloc[:,1].values
```

```
[6] from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,random_state=0)
```

```
print(x.shape)#for 100% of x
print(x_train.shape)#for 75% of x
print(x_test.shape)#for 25% of x
```

```
(30, 1)
(22, 1)
(8, 1)
```

```
[8] from sklearn.linear_model import LogisticRegression
model=LogisticRegression()
```

```
[9] # to train the model, we want both input training and output training data
model.fit(x_train,y_train)
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py:818: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,
LogisticRegression()
```

```
[10] y_pred = model.predict(x_test) #by using the input values ,we predict the output
y_pred
```

```
array([ 39343., 121872., 56957., 57189., 121872., 121872., 121872.,
        57189.])
```

```
[11] y_test#output
```

```
array([ 37731., 122391., 57081., 63218., 116969., 109431., 112635.,
        55794.])
```

```
[12] #individual prediction
model.predict([[1.1]])
```

```
array([39343.])
```

```
[13] model.predict([[4.0]])
```

```
array([57189.])
```

```
[14] model.predict([[4.3]])
```

```
array([56957.])
```

```
[15] model.predict([[4.5]])
```

```
array([61111.])
```

```
[16] df['YearsExperience'].nunique()
```

```
28
```

```
[17] fsalary=df['YearsExperience'].unique()
fsalary
```

```
array([ 1.1, 1.3, 1.5, 2. , 2.2, 2.9, 3. , 3.2, 3.7, 3.9, 4. ,
        4.1, 4.5, 4.9, 5.1, 5.3, 5.9, 6. , 6.8, 7.1, 7.9, 8.2,
        8.7, 9. , 9.5, 9.6, 10.3, 10.5])
```

```
fsize=df.groupby('YearsExperience').size()
fsize
```

```
YearsExperience
1.1    1
1.3    1
1.5    1
2.0    1
2.2    1
2.9    1
3.0    1
3.2    2
3.7    1
3.9    1
4.0    2
4.1    1
4.5    1
4.9    1
5.1    1
5.3    1
5.9    1
6.0    1
```

```


5.3 1
5.9 1
6.0 1
6.8 1
7.1 1
7.9 1
8.2 1
8.7 1
9.0 1
9.5 1
9.6 1
10.3 1
10.5 1
dtype: int64

```

```

[25] import matplotlib.pyplot as plt
plt.bar(fsalary,fsize,color=['cyan','fuchsia','teal','paleturquoise'])


```



```

plt.scatter(df['YearsExperience'],df['Salary'])
plt.title('Scatter Plot')
plt.xlabel('Years of Experience')
plt.ylabel('Salary')

```



```

[27] x=df.iloc[:,0:1].values #input values
x

```

```

array([[ 1.1],
       [ 1.3],
       [ 1.5],
       [ 2. ],
       [ 2.2],
       [ 2.9],
       [ 3. ],
       [ 3.2],
       [ 3.2],
       [ 3.7],
       [ 3.9],
       [ 4. ],
       [ 4. ],
       [ 4.1],
       [ 4.5],
       [ 4.9],
       [ 5.1],
       [ 5.3],
       [ 5.9],
       [ 6. ],
       [ 6.8],
       [ 7.1],
       [ 7.9],
       [ 8.2],
       [ 8.7],
       [ 9. ],
       [ 9.5],
       [ 9.6],
       [10.3],
       [10.5]])

```

```

[28] y=df.iloc[:,1].values #Actual output values
y
array([ 39343., 46205., 37731., 43525., 39891., 56642., 60150.,
       54445., 64445., 57189., 63218., 55794., 56957., 57081.,
       61111., 67938., 66029., 83088., 81363., 93940., 91738.,
       98273., 101302., 113812., 109431., 105582., 116969., 112635.,
       122391., 121872.])

[30] from sklearn.linear_model import LinearRegression #importing linear regression from sklearn.linear_model to perform regression
model = LinearRegression()

[31] #Fit the model -> mapping of inputs with outputs
model.fit(x,y)

LinearRegression()

[32] #predict the model
#using input values we predict the output
y_pred=model.predict(x)
y_pred
array([ 36187.15875227, 38077.15121656, 39967.14368085, 44692.12484158,
       46582.11730587, 53197.09093089, 54142.08716303, 56032.07962732,
       56032.07962732, 60757.06078805, 62647.05325234, 63592.04948449,
       63592.04948449, 64537.04571663, 68317.03064522, 72097.0155738 ,
       73987.00803809, 75877.00050238, 81546.97789525, 82491.9741274 ,
       90051.94398456, 92886.932681 , 100446.90253816, 103281.8912346 ,
       108006.87239533, 110841.86109176, 115566.84225249, 116511.83848464,
       123126.81210966, 125016.80457395])

```

```

[33] y #the actual output values
array([ 39343., 46205., 37731., 43525., 39891., 56642., 60150.,
       54445., 64445., 57189., 63218., 55794., 56957., 57081.,
       61111., 67938., 66029., 83088., 81363., 93940., 91738.,
       98273., 101302., 113812., 109431., 105582., 116969., 112635.,
       122391., 121872.])

[34] #there is difference between our y and y_pred this does not mean our data is inaccurate, it means that our model is not linear
#Linearity of model depends upon size of data and nature of data

[35] #Individual prediction
#We want to predict the salary of 7 years of experience

[36] #model.predict([[x]]),where x=7
model.predict([[7]])
array([91941.93644885])

[40] #By using cross verification technique, y=m*x+c
#to find m and c using python
m=model.coef_
c=model.intercept_
m
array([9449.96232146])

```

```


[41] #to find c using python
c=model.intercept_
c
25792.200198668696

[42] m*7+c
array([91941.93644885])

[43] #The output value of Model prediction and the Cross verification are same,so our Model is 100% Accurate

[44] #Visualization of BEST FIT line
plt.scatter(x,y) #Actual data(points)
plt.plot(x,y_pred,c = 'turquoise') #Predicted data(line)

```



TEAM MEMBERS:

1.Aishwarya R

2.Asad Pangarkar

3.Sai kiran

4.Sarumathe Jayakumar

5.Shailendra Singh

6.Vivek Mangale